

Computer Vision: Project Description

2023-11-30

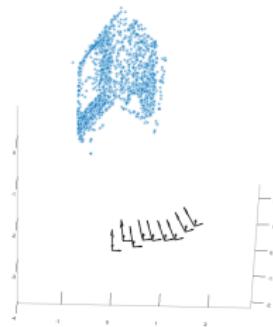
Overview

You should implement a 3D reconstruction software.

In short: from



to



Algorithm: Nonsequential Approach

- ① Calculate relative orientations ($\mathbf{R}_{i,i+1} \mid \mathbf{T}_{i,i+1}$) between images i and $i + 1$
- ② Upgrade to absolute rotations \mathbf{R}_i
- ③ Reconstruct initial 3D points from an initial image pair $i_1 \& i_2$
- ④ For each image i robustly calculate the camera center \mathbf{C}_i / translation \mathbf{T}_i
Reduced camera resectioning problem (analogous to point triangulation)
- ⑤ Refine camera centers (or translation vectors) using Levenberg-Marquardt method
- ⑥ Triangulate points for all pairs $(i, i + 1)$ and visualize 3D points + cameras

Algorithm: Nonsequential Approach

- ① Calculate relative orientations ($R_{i,i+1} | T_{i,i+1}$) between images i and $i + 1$
 - $P_{i,1} = (I | \mathbf{0})$ & $P_{i,2} = (R_{i,i+1} | T_{i,i+1})$
 - Since some datasets have a dominant planar structure:
Implement robust estimation of $(R | T)$ that works for (nearly) planar and general non-planar scenes (following slides) and keep the inliers!
- ② Upgrade to absolute rotations R_i
- ③ Reconstruct initial 3D points from an initial image pair i_1 & i_2
- ④ For each image i *robustly* calculate the camera center C_i / translation T_i
Reduced camera resectioning problem (analogous to point triangulation)
- ⑤ Refine camera centers (or translation vectors) using Levenberg-Marquardt method
- ⑥ Triangulate points for all pairs $(i, i + 1)$ and visualize 3D points + cameras

Algorithm: Nonsequential Approach

- ① Calculate relative orientations ($\mathbf{R}_{i,i+1} \mid \mathbf{T}_{i,i+1}$) between images i and $i + 1$
- ② Upgrade to absolute rotations \mathbf{R}_i
 - Set $\mathbf{R}_1 = \mathbf{I}$ and chain rotation matrices $\mathbf{R}_i = \mathbf{R}_{i-1,i}\mathbf{R}_{i-1}$
 - You now know the rotations matrices for all cameras!
- ③ Reconstruct initial 3D points from an initial image pair $i_1 \& i_2$
- ④ For each image i *robustly* calculate the camera center \mathbf{C}_i / translation \mathbf{T}_i
Reduced camera resectioning problem (analogous to point triangulation)
- ⑤ Refine camera centers (or translation vectors) using Levenberg-Marquardt method
- ⑥ Triangulate points for all pairs $(i, i + 1)$ and visualize 3D points + cameras

Algorithm: Nonsequential Approach

- ① Calculate relative orientations ($R_{i,i+1} | T_{i,i+1}$) between images i and $i + 1$
- ② Upgrade to absolute rotations R_i
- ③ Reconstruct initial 3D points from an initial image pair i_1 & i_2
 - For better accuracy: i_1 and i_2 should not be adjacent (have enough baseline)
Suggested values provided by us
 - Calculate ($R_{i_2,i_1} | T_{i_2,i_1}$) and triangulate inliers: 3D points \mathcal{X}_0
 - Center \mathcal{X}_0 and bring to world coordinates (via $R_{i_1}^\top$)
 - Save the SIFT descriptors (coming from i_1 or i_2) for reconstructed 3D points
e.g. in desc_X
- ④ For each image i robustly calculate the camera center C_i / translation T_i
Reduced camera resectioning problem (analogous to point triangulation)
- ⑤ Refine camera centers (or translation vectors) using Levenberg-Marquardt method
- ⑥ Triangulate points for all pairs $(i, i + 1)$ and visualize 3D points + cameras

Algorithm: Nonsequential Approach

- ① Calculate relative orientations ($R_{i,i+1} | T_{i,i+1}$) between images i and $i + 1$
- ② Upgrade to absolute rotations R_i
- ③ Reconstruct initial 3D points from an initial image pair i_1 & i_2
- ④ For each image i robustly calculate the camera center C_i / translation T_i
Reduced camera resectioning problem (analogous to point triangulation)
 - ① Establish correspondences between i and 3d points (using desc_X), e.g.
`[matches_2d_3d, scores] = vl_ubcmatch(descriptors{i}, desc_X);`
 - ② Estimate C_i/T_i robustly from these correspondences using a 2-point method
- ⑤ Refine camera centers (or translation vectors) using Levenberg-Marquardt method
 - Modify CE in assignment 4
- ⑥ Triangulate points for all pairs $(i, i + 1)$ and visualize 3D points + cameras

Things you need to implement

In order to complete the mandatory parts

- ① A RANSAC algorithm robustly estimates \mathbf{R} and \mathbf{T} from correspondences
 - For scenes with and w/o a dominant plane
 - Suggested approach: use a “parallel” version of RANSAC
 - Alternative: implement a direct 5-point solver (optional points)
 - In this case a standard RANSAC loop is sufficient
- ② Extract \mathbf{R} and \mathbf{T} from either \mathbf{E} or \mathbf{H}
 - $(\mathbf{R}|\mathbf{T})$ from \mathbf{E} : CE3 in assignment 3
 - $(\mathbf{R}|\mathbf{T})$ from \mathbf{H} : provided (`homography_to_RT`)
 - Use cheirality to determine the correct configuration
- ③ A method to estimate the translation \mathbf{T} robustly from 2D-3D correspondences $\mathbf{x}_i \leftrightarrow \mathbf{X}_i$
 - E.g. `P = estimate_T_robust(xs, Xs, inlier_threshold)`
 - Simplify `estimate_camera_DLT` from CE2 in assignment 2 to estimate \mathbf{T} for a minimal sample
- ④ High-level wrapper code (“glue code”)

RANSAC: Searching for E and H in parallel

Modify the vanilla RANSAC loop

- Maintain (and return) $(R^*|T^*)$
- Maintain inlier fractions ε_E and ε_H and corresponding number of iterations
- Return best $(R^*|T^*)$ found after sufficient iterations for a good E or H

In each iteration

- Estimate E using the 8/7-point method from a minimal sample set of size 8/7
 - If E is the new best essential matrix:
Adjust ε_E , number of E-iterations, and extract new $(R^*|T^*)$ from E
- Estimate H using the DLT from a minimal sample set of size 4
 - If H is new best homography:
Extract two solutions $(R_a|T_a)$ and $(R_b|T_b)$ from H (homography_to_RT)
Test both essential matrices $E \in \{[T_a] \times R_a, [T_b] \times R_b\}$ for validity
If E is the new best essential matrix and is "acceptable"
Assign correct $(R|T)$ (extracted from E) to $(R^*|T^*)$
Adjust ε_H & number of H-iterations
Adjust ε_E & number of E-iterations

An essential matrix E is "acceptable" if exactly one of the 4 cheirality configurations has all inliers (w.r.t. the epipolar constraint) in front of both cameras

RANSAC: Searching for E and H in parallel

Modify the vanilla RANSAC loop

- Maintain (and return) $(R^*|T^*)$
- Maintain inlier fractions ε_E and ε_H and corresponding number of iterations
- Return best $(R^*|T^*)$ found after sufficient iterations for a good E or H

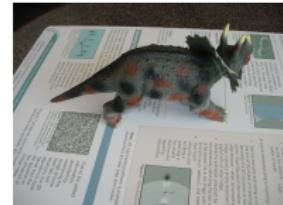
In each iteration

- Estimate E using the 8/7-point method from a minimal sample set of size 8/7
 - If E is the new best essential matrix:
Adjust ε_E , number of E-iterations, and extract new $(R^*|T^*)$ from E
- Estimate H using the DLT from a minimal sample set of size 4
 - If H is new best homography:
Extract two solutions $(R_a|T_a)$ and $(R_b|T_b)$ from H (homography_to_RT)
Test both essential matrices $E \in \{[T_a] \times R_a, [T_b] \times R_b\}$ for validity
If E is the new best essential matrix *and* is "acceptable"
Assign correct $(R|T)$ (extracted from E) to $(R^*|T^*)$
Adjust ε_H & number of H-iterations
Adjust ε_E & number of E-iterations

An essential matrix E is "acceptable" if exactly one of the 4 cheirality configurations has all inliers (w.r.t. the epipolar constraint) in front of both cameras

How to run your software

- Once your implementation is complete, you should run it (in Matlab) via
> `run_sfm(<dataset number>)`
- Datasets 1–9 are test datasets provided by us
 - Datasets 1 & 2: for debugging, do not include in the report
 - Datasets 3–7: increasing difficulty, you should include those in the report
 - Datasets 8 & 9: optional
- You may include your own datasets (optional, see later)



For the report...

- Provide a high-level explanation of your algorithmic design choices
 - Especially what you implemented beyond the routines from the computer exercises
 - Also if you modified/improved your implementation for the computer exercises
- Describe unexpected problems you encountered (if any)
- Document changes to the provided code
 - E.g. modifying inlier threshold or selecting a different initial pair
- Briefly describe extra datasets (if you have any)
 - If you picked something from the internet: link to the dataset(s)
 - Describe unexpected issues with the additional datasets
- For optional points: describe what you implemented

Deadline: Thursday, January 4nd, 2024, 23:59

Extracting R and T from H

- Bill Triggs, “Autocalibration from Planar Scenes”, ECCV 1998 (Appendix 1)
 - <http://perception.inrialpes.fr/Publications/1998/Tri98/Triggs-eccv98-long.pdf>
 - Adapted to our conventions
 - function [R1, t1, R2, t2] = homography_to_RT(H, x1, x2)
 - Python: convert Matlab routine, or
 - num, Rs, Ts, Ns = cv2.decomposeHomographyMat(H, K) (with K = I)

```
function [R1,t1,n1, R2,t2,n2, zeta] = homog_to_Rt(H)
[U,S,V] = svd(H);
s1 = S(1,1)/S(2,2); s3 = S(3,3)/S(2,2);
zeta = s1-s3;
a1 = sqrt(1-s3^2); b1 = sqrt(s1^2-1);
[a,b] = unitize(a1,b1);
[c,d] = unitize( 1+s1*s3, a1*b1 );
[e,f] = unitize( -b/s1, -a/s3 );
v1 = V(:,1); v3 = V(:,3);
n1 = b*v1-a*v3; n2 = b*v1+a*v3;
R1 = U*[c,0,d; 0,1,0; -d,0,c]*V';
R2 = U*[c,0,-d; 0,1,0; d,0,c]*V';
t1 = e*v1+f*v3; t2 = e*v1-f*v3;
if (n1(3)<0) t1 = -t1; n1 = -n1; end;
if (n2(3)<0) t2 = -t2; n2 = -n2; end;
end;
```

Importing dataset information

- You may create a function `get_dataset_info(<dataset number>)`, e.g.

```
function [K, image_names, init_pair] = get_dataset_info(dataset)
    if dataset == 1
        image_names = {"dataset1/kronan1.JPG", "dataset1/kronan2.JPG"};
        im_width = 1936; im_height = 1296;
        focal_length_35mm = 45.0; % from the EXIF data
        init_pair = [1, 2];
        pixel_threshold = 1.0;
    elseif
        ...
    end
    focal_length = max(im_width, im_height) * focal_length_35mm / 35.0
    K = [focal_length 0 im_width/2; 0 focal_length im_height/2; 0 0 1]
end
```

- This is probably one of the first routines to call in your `run_sfm` script
 - Another one is setting up `vl_feat`
- *Provided together with the project datasets in the data folder*

Practical Suggestions

- Work with normalized image points $\tilde{\mathbf{x}}_i = \mathbf{K}^{-1} \mathbf{x}_i$
 - Camera matrices are $\mathbf{P}_i = (\mathbf{R}_i \mid \mathbf{T}_i)$
 - Do not forget to adjust inlier thresholds from pixels to normalized units
(`pixel_threshold` is a suggested value in `get_dataset_info`)

```
epipolar_threshold      = pixel_threshold / focal_length
homography_threshold   = 3 * pixel_threshold / focal_length
translation_threshold  = 3 * pixel_threshold / focal_length
```
 - Use maybe `xn` for normalized and `xu` for unnormalized image points
- Filter 3D points excessively far away from the center of gravity, e.g.

$$\|\mathbf{X}_j - \bar{\mathbf{X}}\| \leq 5 \times 90\% \text{ quantile of } \{\|\mathbf{X}_j - \bar{\mathbf{X}}\|\}_j$$

- Check (visualize) triangulated points and cameras for each relative pose
 - $\mathbf{P}_{i,1} = (\mathbf{I} \mid \mathbf{0})$ & $\mathbf{P}_{i,2} = (\mathbf{R}_{i+1,i} \mid \mathbf{T}_{i+1,i})$
 - Verify correctness of relative pose and cheirality
- While in the development/debugging phase
 - Precompute and store keypoints and matches
 - Load those for faster time to reach your own code

Practical Suggestions

- Work first with the non-planar datasets 1-5 and start the RANSAC implementation from CE2 in assignment 4
- RANSAC is a stochastic method, so it may *occasionally* fail
 - But there is likely a bug if your software frequently produces bad 3D models
- Use `randperm(N, K)` to draw K indices from $\{1, \dots, N\}$ (w/o replacement)
- Reproducability while debugging: early in your code set the random seed via
 - Matlab: `rng(42);`
 - Octave: `rand("state", 42);`
- Accelerating model evaluation in RANSAC
 - Use `.^2`, `sum(X, dim)` and `sqrt.` instead of loops
E.g. `sqrt.(sum((x-y).^2, 1))`

Optional Points for Additional Functionality

- ① Run your software on datasets 8 & 9 (5 points each, max. 10 points)
- ② Run your software on more datasets (10 points/dataset, max. 20 points)
 - Describe dataset (and its origin) and your experiences briefly
 - Recall lecture 3 for how to extract K
 - Images should have \approx HD (1920×1080) resolution
 - Tradeoff between run-time and image details
 - Use dataset ids ≥ 10
- ③ Use MSAC model scoring (max. 10 points)
 - Reference: <https://www.robots.ox.ac.uk/~vgg/publications/2000/Torr00/torr00.pdf>
- ④ Faster RANSAC using early bailout: $T_{d,d}$ test (max. 10 points)
 - Reference: Chum & Matas, "Randomized RANSAC with $T_{d,d}$ test"
- ⑤ Nonlinear refinement of camera pose (max. 20 points)
 - Optimize over camera translation/center and rotation matrix
- ⑥ Implement the 7-point method (cf. lecture 8, max. 20 points)
 - Document how you used a CAS
- ⑦ Implement the 5-point method (max. 60 points)
 - Cf. lecture 8 and the following hints
 - No need for the parallel search for E and H
 - Document how you used a CAS (if you did)

The 5-Point Method

Finding an essential matrix can be done with five correspondences by solving

$$\bar{\mathbf{x}}_i^T \mathbf{E} \mathbf{x}_i = 0 \quad i = 1, \dots, 5$$

$$\det(\mathbf{E}) = 0$$

$$2\mathbf{E}\mathbf{E}^\top - \text{trace}(\mathbf{E}\mathbf{E}^\top)\mathbf{E} = 0$$

Form the system matrix \mathbf{M} from the five correspondences. \mathbf{M} has a 4 dimensional null-space:

$$\mathbf{M}(x\mathbf{n}_1 + y\mathbf{n}_2 + z\mathbf{n}_3 + w\mathbf{n}_4) = 0$$

Reshape into matrices \mathbf{E}_i such that $\mathbf{n}_i = \text{vec}(\mathbf{E}_i)$

$$\bar{\mathbf{x}}_i^T (x\mathbf{E}_1 + y\mathbf{E}_2 + z\mathbf{E}_3 + w\mathbf{E}_4) \mathbf{x}_i = 0 \quad i = 1, \dots, 5$$

The 5-Point Method

- Use the remaining equations to determine x, y, z, w
 - Fix $w = 1$: $\mathbf{E} = x\mathbf{E}_1 + y\mathbf{E}_2 + z\mathbf{E}_3 + \mathbf{E}_4$
- Identify $x \rightarrow \alpha_1, y \rightarrow \alpha_2, z \rightarrow \alpha_3, 1 \rightarrow \alpha_4$
Trace constraint, 9 equations

$$2\mathbf{EE}^\top \mathbf{E} - \text{trace}(\mathbf{EE}^\top) \mathbf{E} = \sum_{i=1}^4 \sum_{j=1}^4 \sum_{k=1}^4 \alpha_i \alpha_j \alpha_k \left(2\mathbf{E}_i \mathbf{E}_j^\top \mathbf{E}_k - \text{trace}(\mathbf{E}_i \mathbf{E}_j^\top) \mathbf{E}_k \right) \stackrel{!}{=} 0$$

Determinant, one equation:

$$\det(\mathbf{E}) = \sum_{i=1}^4 \sum_{j=1}^4 \sum_{k=1}^4 \alpha_i \alpha_j \alpha_k (e_{11}^i e_{22}^j e_{33}^k + e_{12}^i e_{23}^j e_{31}^k + e_{13}^i e_{21}^j e_{32}^k - e_{11}^i e_{23}^j e_{32}^k - e_{12}^i e_{21}^j e_{33}^k - e_{13}^i e_{22}^j e_{31}^k) \stackrel{!}{=} 0$$

- Map 64 expressions $\alpha_i \alpha_j \alpha_k$ to the 20 degree ≤ 3 monomials in x, y, z
 - E.g. $\alpha_1 \alpha_2 \alpha_1 \rightarrow x^2 y$ (index 2) and $\alpha_3 \alpha_4 \alpha_4 \rightarrow z$ (index 19)
 - You can fill a 3-tensor $T(i, j, k)$ to map (i, j, k) to its monomial index
- Or use a CAS to extract the coefficients for the 10 multi-variate polynomials
 - Analogous to the use of a CAS for the 7-point method (lecture 8)
- You obtain 10×20 coefficients!

The 5-Point Method

- Use the remaining equations to determine x, y, z, w
 - Fix $w = 1$: $\mathbf{E} = x\mathbf{E}_1 + y\mathbf{E}_2 + z\mathbf{E}_3 + \mathbf{E}_4$
- Identify $x \rightarrow \alpha_1, y \rightarrow \alpha_2, z \rightarrow \alpha_3, 1 \rightarrow \alpha_4$
Trace constraint, 9 equations

$$2\mathbf{E}\mathbf{E}^\top \mathbf{E} - \text{trace}(\mathbf{E}\mathbf{E}^\top)\mathbf{E} = \sum_{i=1}^4 \sum_{j=1}^4 \sum_{k=1}^4 \alpha_i \alpha_j \alpha_k \left(2\mathbf{E}_i \mathbf{E}_j^\top \mathbf{E}_k - \text{trace}(\mathbf{E}_i \mathbf{E}_j^\top) \mathbf{E}_k \right) \stackrel{!}{=} 0$$

Determinant, one equation:

$$\det(\mathbf{E}) = \sum_{i=1}^4 \sum_{j=1}^4 \sum_{k=1}^4 \alpha_i \alpha_j \alpha_k (e_{11}^i e_{22}^j e_{33}^k + e_{12}^i e_{23}^j e_{31}^k + e_{13}^i e_{21}^j e_{32}^k - e_{11}^i e_{23}^j e_{32}^k - e_{12}^i e_{21}^j e_{33}^k - e_{13}^i e_{22}^j e_{31}^k) \stackrel{!}{=} 0$$

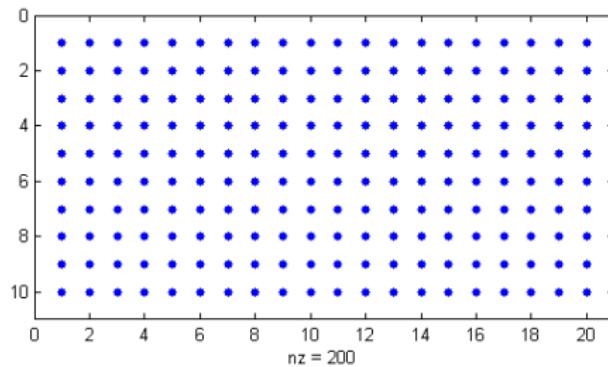
- Map 64 expressions $\alpha_i \alpha_j \alpha_k$ to the 20 degree ≤ 3 monomials in x, y, z
 - E.g. $\alpha_1 \alpha_2 \alpha_1 \rightarrow x^2 y$ (index 2) and $\alpha_3 \alpha_4 \alpha_4 \rightarrow z$ (index 19)
 - You can fill a 3-tensor $T(i, j, k)$ to map (i, j, k) to its monomial index
- Or use a CAS to extract the coefficients for the 10 multi-variate polynomials
 - Analogous to the use of a CAS for the 7-point method (lecture 8)
- You obtain 10×20 coefficients!

The 5-Point Method

- Order for monomials: all 10 degree-3 monomials come first, e.g.
 - Correspond to columns in your 10×20 coefficient matrix C

$$(x^3, x^2y, x^2z, xy^2, xyz, xz^2, y^3, y^2z, yz^2, z^3, \quad x^2, xy, xz, y^2, yz, z^2, x, y, z, 1)$$

- Gaussian elimination gives reductions for all third order terms
 - Lucky coincidence: 10 degree-3 monomials and 10 equations
 - Gaussian elimination: $C' = (C_{1:10, 1:10})^{-1} C$

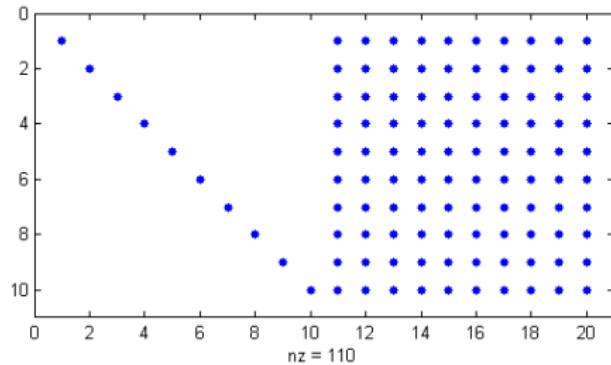


The 5-Point Method

- Order for monomials: all 10 degree-3 monomials come first, e.g.
 - Correspond to columns in your 10×20 coefficient matrix C

$$(x^3, x^2y, x^2z, xy^2, xyz, xz^2, y^3, y^2z, yz^2, z^3, \quad x^2, xy, xz, y^2, yz, z^2, x, y, z, 1)$$

- Gaussian elimination gives reductions for all third order terms
 - Lucky coincidence: 10 degree-3 monomials and 10 equations
 - Gaussian elimination: $C' = (C_{1:10, 1:10})^{-1} C$



The 5-Point Method

- We can express all degree-3 monomials via degree- ≤ 2 ones
- Action matrix M_x (or M_y etc) is 10×10 matrix
 - Only 6 equations / rows of C' will be used

$$x \begin{pmatrix} x^2 \\ xy \\ xz \\ y^2 \\ yz \\ z^2 \\ x \\ y \\ z \\ 1 \end{pmatrix} = \begin{pmatrix} x^3 \\ x^2y \\ x^2z \\ xy^2 \\ xyz \\ xz^2 \\ x^2 \\ xy \\ xz \\ x \end{pmatrix} = M_x \begin{pmatrix} x^2 \\ xy \\ xz \\ y^2 \\ yz \\ z^2 \\ x \\ y \\ z \\ 1 \end{pmatrix}$$

- Up to 10 real eigenvalues for x
- How do we get y and z ?
 - Each eigenvector $v \sim (x^2, xy, xz, y^2, yz, z^2, x, y, z, 1)^\top$
 - Extract (x, y, z) as $v_{7:9}/v_{10}$
 - Discard complex solutions (what about “almost” real ones?)
 - Or try the real parts of all (x, y, z)
 - $E = xE_1 + yE_2 + zE_3 + E_4$
 - Evaluate all 10 essential matrices for their inliers

The 5-Point Method

- We can express all degree-3 monomials via degree- ≤ 2 ones
- Action matrix M_x (or M_y etc) is 10×10 matrix
 - Only 6 equations / rows of C' will be used

$$x \begin{pmatrix} x^2 \\ xy \\ xz \\ y^2 \\ yz \\ z^2 \\ x \\ y \\ z \\ 1 \end{pmatrix} = \begin{pmatrix} x^3 \\ x^2y \\ x^2z \\ xy^2 \\ xyz \\ xz^2 \\ x^2 \\ xy \\ xz \\ x \end{pmatrix} = M_x \begin{pmatrix} x^2 \\ xy \\ xz \\ y^2 \\ yz \\ z^2 \\ x \\ y \\ z \\ 1 \end{pmatrix}$$

- Up to 10 real eigenvalues for x
- How do we get y and z ?
- Each eigenvector $v \sim (x^2, xy, xz, y^2, yz, z^2, x, y, z, 1)^\top$
- Extract (x, y, z) as $v_{7:9}/v_{10}$
 - Discard complex solutions (what about “almost” real ones?)
 - Or try the real parts of all (x, y, z)
 - $E = xE_1 + yE_2 + zE_3 + E_4$
- Evaluate all 10 essential matrices for their inliers